

TitleTrack

Pick your city job title, see what AI actually means for your work task-by-task — no fake risk scores — and see the exam that's your next move up.

Ryan Michie

Built on NYC's own payroll, exam, and occupation data

“How do we make sure AI lifts workers instead of leaving them behind?”

— the question this event asked

For 550,219 people on New York City’s payroll, that question is currently answered by headlines and generic “40% of jobs at risk” scores. Those scores are wrong, because they treat a job title as if it were one thing. The city’s own data proves it isn’t:

- **CARETAKER** is a NYCHA janitor, not a caregiver.
- **TEACHER- PER SESSION** — 78,618 people, the largest title in the city — isn’t a job at all. It’s a pay code.
- **LIEUTENANT** is two different jobs, NYPD or FDNY, mapping to two different occupation codes.
- **And nothing connects.** Payroll, exams, and occupations share no key. The join has to be rebuilt from text.

Answer it one task at a time. Never score the job.

A worker picks their real payroll title and sees the real tasks that job involves — published federal task statements, shown verbatim. Every statement gets one of three labels:

AUTOMATABLE

AUGMENTABLE

HUMAN ANCHORED

- **No overall risk score — on purpose.** Holding method constant and changing only the unit of analysis moves US “high-risk” employment from 38% to 9% (Arntz et al. 2017). The scary number is an artifact of aggregation.
- **Two kinds of human, and they are opposites.** A task can stay human because it needs a body present — which erodes as technology improves — or because the law vests the act in a person, which no model release changes. A technology roadmap and a policy roadmap. One score erases the difference.
- **It knows when to refuse.** Ask it about CITY SEASONAL AIDE and it declines: a hiring category, 17 salaried rows out of 2,414. A gated eval fails the build if that refusal is ever replaced with a guess.
- **Then it gives them a move.** Open DCAS exams for their title, plus their real promotional next rung — the actual civil-service mechanism for moving up.

I built the test first, then let it fail.

- **Wrote the test before the app.** It rebuilds every number on the page from the city's own files. If anything doesn't match, the build fails.
- **Agent 1 — title-matcher.** Matches exam names to payroll names. It is not allowed to force a match. Guessing produced 12 wrong ones.
- **Agent 2 — crosswalk-decider.** Picks the job code for a title using evidence in the files. It says BLOCKED instead of guessing. It corrected me twice.
- **Two questions, never mixed.** Can AI do this task? And is a person legally required to be the one who does it? The second is about law, not technology.
- **Claude explains; code does the math.** Python handles the data. Claude only puts it in plain English — with a saved answer ready if the internet drops.

Helping AI lift workers, not leave them behind.

- **Workers can see what AI means for their job.** Task by task, for their real title — traffic agents, school safety agents, EMTs, caretakers.
- **They get a real next step.** The city already has a promotion ladder. But nobody tells you when your exam opens, and the window is only three weeks.
- **Training goes where it pays off.** Learn the tasks AI will help with. Skip the ones the law says only a person can do.
- **The results surprise people.** The so-called low-skill jobs are the least exposed. A school secretary is more exposed than a NYCHA caretaker.
- **Next: ask the workers if we got it right.** Test the labels with union members, and pilot exam alerts with the city.