CLAUDE COMMUNITY IMPACT LAB · PROJECT STATUS · AUGUST 2026

TitleTrack

What can AI actually do — and be trusted to do — for New York City's public workforce?

The problem

NYC pays 550,219 people under 1,557 job titles that routinely hide the real work — the city's single largest "title" is not a job at all. Agencies are being sold AI-readiness scores that no one can check against reality.

The solution

An open, evidence-first crosswalk — payroll titles → civil-service exams → O*NET occupations — plus a rubric that separates "a model can produce this output" from "a model can be trusted with it unsupervised."

The community value

A defensible public map of where AI genuinely helps city work — and where authority must stay with an accountable human. In public agencies those differ, and naming that gap honestly is what protects New Yorkers.

Where the project stands today

Findings that protect NYC from bad AI math

78,618 people sit under the city's largest payroll "title" — a pay code, not a job. Coding it as work double-counts the teaching workforce.

30.1% of top-40 headcount is not an occupation at all (poll workers, hiring categories). Excluding it is accuracy — a finding, not a gap.

13 / 40 hardest title→occupation calls decided, evidence on record — including 3 errors in our own docs, found and fixed by our own checks.

Foundation already built: 6.8M payroll rows · 2,901 exams · O*NET 30.3 occupations · byte-reproducible data builds on any machine · scoring rubric v0.1 with 5 hand-scored tasks and its ambiguities documented

Two custom Claude agents — each with a spec, golden tests, hard gates

`title-matcher`

Matches exam titles to payroll titles without ever forcing one. 54 golden cases drawn from the real join residuals. The fuzzy baseline scores 37/41 recall — and FAILS with 12 corrupt matches. Wrong answers fail the run; misses only cost recall.

`crosswalk-decider`

Turns titles into occupation codes with evidence ranked above word-similarity. Graded against our 13 human decisions, and required to answer BLOCKED when the deciding fact is not in the data. Baselines fail in both directions — by design.

The road ahead

1. **Unblock** — *days*
   One live API query (agency × title) closes LIEUTENANT and confirms two more calls. Needs any laptop — the sandbox blocks NYC Open Data.

2. **Decide** — *weeks 1–2*
   Run crosswalk-decider on the 27 open titles; a human signs off on every verdict.

3. **Review** — *weeks 2–3*
   A civil-service insider checks our title decodings; counsel reviews the rubric’s legal “authority” tests.

4. **Match** — *weeks 3–4*
   title-matcher lifts occupation coverage past 32.3% of headcount — with zero forced matches.

5. **Deliver** — *weeks 5+*
   Rubric v0.2, task scoring at volume, and an agency-facing report with a live demo.

Decisions the team owns (not Claude): how to classify task statements that name no consequence (rubric ambiguity A2) · the grade-level roll-up affecting 101,630 teachers · what shape the public deliverable takes

How a small team runs this with Claude

One source of truth

The GitHub repo. CLAUDE.md onboards every teammate — and every Claude session — with the same data hazards and conventions. Agents, evals, and raw snapshots live in-repo, so nothing is tribal knowledge.

No GitHub account? Use a Claude Project

Create a shared Claude Project and upload CLAUDE.md, the decision sheet, and the rubric as project knowledge. Teammates review evidence and make calls in chat; one repo-holder commits the outcomes back.

Delegate by artifact, not by tool

Decision sheet, rubric, matcher — each workstream has its own oracle: a gated eval or a byte-identical build check. Work gets graded before it merges, whoever (or whatever) produced it.

Let the gates referee

Evals hard-fail wrong answers, never misses — the same rule for humans and agents. Nobody edits golden cases to pass; if you disagree with a label, bring evidence. That route has already fixed three labels.

Definition of done, every week: all evals PASS · data build byte-identical (--check) · every correction logged in-repo, not in someone's head